


Container build at scale using BuildKit and BuildX



About Me

Lead AI Ops SRE @ Cisco

17 years in the IT (Ex-SUSE/Rancher/Nutanix/Oracle/HPE/IBM)

Blog: <https://ansilh.com>



Agenda

- Introduction to Docker build
- BuildKit vs. Legacy Docker Build
- BuildKit LLB
- BuildKit Drivers (Docker, Kubernetes, Remote)
- BuildKit Security
- Multi-Platform Builds
- Bake HCL Workflows
- Cache
- Demos



Docker build - Legacy

- Sequential execution
- Each instruction generate a cached layer
- Caching relies on heuristics (e.g., timestamps, Dockerfile instruction text, or file metadata)
- Secrets passed via ARG have security risks
- Multi-platform builds needs workarounds



Docker build - BuildKit

- Graph based execution using Low Level Build (LLB)
- Executes independent steps in parallel
- Supports remote caching and inline caching
- Secrets can be securely mounted using `--mount=type=secret`
- Rootless mode allows builds without root privilege
- Native support for different architectural builds via QEMU or remote builders




BuildKit vs. Legacy Docker Build

Legacy Build	BuildKit
Sequential stages	Parallel stage execution
Basic local cache	Multi-layer caching
Limited platform support	Native multi-platform
No built-in secrets	<code>--mount=type=secret</code>
No frontend customization	Extensible LLB (Low-Level Build)



BuildKit LLB

- Low Level Build format is an intermediate binary format
- LLB defines a content-addressable dependency graph



Frontend

The “shabang” for build definition

<https://docs.docker.com/build/buildkit/frontend/>

Flexibility to get features and bug fixes without docker engine updates

```
# syntax=[remote image reference]
```

```
# syntax=docker/dockerfile:1
```

```
# syntax=docker.io/docker/dockerfile:1
```

```
# syntax=example.com/user/repo:tag@sha256:abcdef...
```



BuildX

- A CLI that adds additional features to docker build
- High-level build constructs like “bake”
- Installation - refer <https://github.com/docker/buildx>



BuildKit Drivers

1. Docker Driver (default, uses Docker's daemon).
2. Docker Container Driver
3. Kubernetes Driver (scale builds across pods).
 - Use case: CI/CD in k8s clusters.
4. Remote Driver (connect to external BuildKitd).
 - Use case: Dedicated BuildKit servers.



Docker driver

The default driver

Built directly into the docker engine

Images built will be loaded directly to the image store

```
ansil@ansil-NUC10i7FNH:~/cncg-kochi/demo-1$ docker buildx ls
NAME/NODE      DRIVER/ENDPOINT   STATUS   BUILDKIT   PLATFORMS
container*     docker-container
 \_ container0 \_ unix:///var/run/docker.sock running   v0.20.1    linux/amd64, linux/amd64/v2, linux/amd64/v3, linux/386
default        docker
 \_ default    \_ default        running   v0.13.1    linux/amd64, linux/amd64/v2, linux/amd64/v3, linux/386
```



Docker Container Driver

- Specify custom BuildKit versions to use.
- Build multi-arch images
- Advanced options for cache import and export

```
docker buildx create \  
  --name container \  
  --driver=docker-container \  
  --driver-opt=image=moby/buildkit:v0.20.1 \  
  --use --bootstrap
```



Docker Container Driver

```
ansil@ansil-NUC10i7FNH:~/cncg-kochi/demo-1$ docker buildx create \
--name container \
--driver=docker-container \
--driver-opt=image=moby/buildkit:v0.20.1 \
--use --bootstrap
[+] Building 3.0s (1/1) FINISHED
=> [internal] booting buildkit
=> => pulling image moby/buildkit:v0.20.1
=> => creating container buildx_buildkit_container0
container
ansil@ansil-NUC10i7FNH:~/cncg-kochi/demo-1$ docker ps
CONTAINER ID   IMAGE                COMMAND                  CREATED        STATUS        PORTS        NAMES
d4cbbd445781   moby/buildkit:v0.20.1 "buildkitd --allow-i..." 3 seconds ago  Up 3 seconds          buildx_buildkit_container0
ansil@ansil-NUC10i7FNH:~/cncg-kochi/demo-1$ docker buildx ls
NAME/NODE      DRIVER/ENDPOINT   STATUS    BUILDKIT   PLATFORMS
container*     docker-container  running   v0.20.1    linux/amd64, linux/amd64/v2, linux/amd64/v3, linux/386
default        docker            running   v0.13.1    linux/amd64, linux/amd64/v2, linux/amd64/v3, linux/386
ansil@ansil-NUC10i7FNH:~/cncg-kochi/demo-1$
```

Diagram illustrating the Docker buildx workflow:

1. `docker buildx create` command is executed.
2. Buildkit is booting and pulling the image `moby/buildkit:v0.20.1`.
3. Container `buildx_buildkit_container0` is created.
4. The container is running and ready for use.

time docker buildx build --platform=linux/amd64,linux/arm64

--builder=container -t ansilh/buildkit-demo:v0 --push --progress=plain .



Docker container driver - multi-platform build

dockerhub Explore My Hub Search Docker Hub

ansilh Docker Personal

Repositories / buildkit-demo / Tags / v0

ansilh/buildkit-demo:v0 MULTI-PLATFORM

INDEX DIGEST sha256:8a4a6d2ebabfc7008ef1288bee7274c6669344970a0d2d4ec983385b73f2f3c

OS/ARCH linux/arm64

COMPRESSED SIZE 1.76 MB

LAST PUSHED 1 minute by ansilh

TYPE Image

MANIFEST DIGEST sha256:5b4051a6...

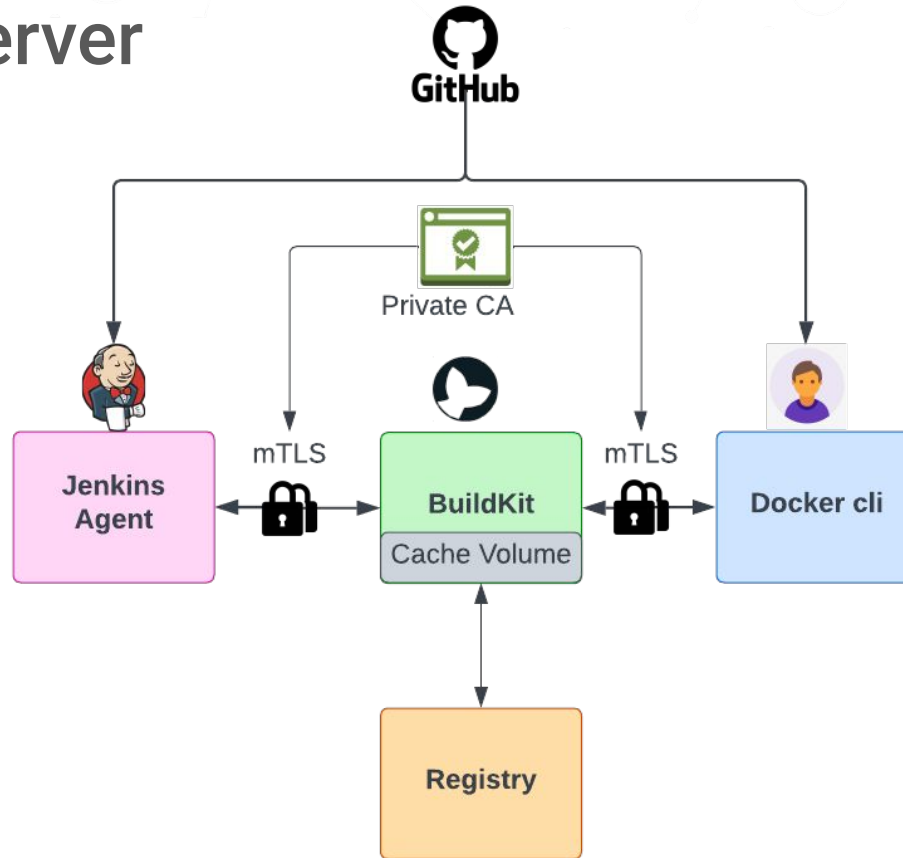
Image Layers Vulnerabilities

Image Layers

	Command
1 BusyBox 1.37.0 (glibc), Debian 12 1.76 MB	BusyBox 1.37.0 (glibc), Debian 12
2 RUN /bin/sh -c echo BEGIN:st2 108 B	
3 COPY /ST0 /ST2.A # buildkit 93 B	
4 COPY /ST1 /ST2.B # buildkit 93 B	
5 RUN /bin/sh -c echo END:st2 32 B	



Buildkit Server



BuildKit Security

Secrets Management:

```
RUN --mount=type=secret,id=my_secret cat /run/secrets/my_secret
```

Rootless Mode: Run builds without root privileges.

<https://github.com/rootless-containers/rootlesskit/>





Bake



kubernetes

Buildx Bake

Bake is a command built into the Buildx CLI,

Bake is an abstraction for the docker build command that lets you more easily manage your build configuration (CLI flags, environment variables, etc.) in a consistent way for everyone on your team.

You can write Bake files in HCL, YAML (Docker Compose files), or JSON. In general, HCL is the most expressive and flexible format



Example

```
# docker build -f Dockerfile -t myapp:latest .
```

docker-bake.hcl

```
target "myapp" {  
    context = "."  
    dockerfile = "Dockerfile"  
    tags = ["myapp:latest"]  
}  
# docker buildx bake myapp
```



Targets

```
# Custom target
target "webapp" {
  dockerfile = "webapp.Dockerfile"
  tags = ["docker.io/ansilh/webapp:latest"]
  context = "https://github.com/ansilh/webapp"
}

# Default target
target "default" {
  dockerfile = "webapp.Dockerfile"
  tags = ["docker.io/ansilh/webapp:latest"]
  context = "https://github.com/username/webapp"
}

# docker buildx bake --print
```



Inheritance

```
target "_common" {
  args = {
    GO_VERSION = "1.23"
    BUILDKIT_CONTEXT_KEEP_GIT_DIR = 1
  }
}

target "app-dev" {
  inherits = ["_common"]

  args = {
    BUILDKIT_CONTEXT_KEEP_GIT_DIR = 0
  }

  tags = ["docker.io/ansilh/myapp:dev"]
  labels = {
    "org.opencontainers.image.source" = "https://github.com/ansilh/myapp"
    "org.opencontainers.image.author" = "ansilh@example.com"
  }
}

target "app-release" {
  inherits = ["app-dev", "_common"]
  tags = ["docker.io/ansilh/myapp:latest"]
  platforms = ["linux/amd64", "linux/arm64"]
}
```



Variables

```
group "default" {  
    targets = [ "webapp" ]  
}
```

```
variable "TAG" {  
    default = "latest"  
}
```

```
target "webapp" {  
    context = "."  
    dockerfile = "Dockerfile"  
    tags = ["docker.io/ansilh/webapp: ${TAG}"]  
}
```



Functions

Bake ships with built-in support for the go-cty

<https://github.com/zclconf/go-cty/blob/main/README.md>

```
variable "TAG" {  
    default = "latest"  
}  
  
group "default" {  
    targets = ["webapp"]  
}  
  
target "webapp" {  
    args = {  
        buildno = "${add(123, 1)}"  
    }  
}
```

<https://github.com/zclconf/go-cty/blob/main/cty/function/stdlib/number.go#L30>




User defined functions

```
function "increment" {  
    params = [number]  
    result = number + 1  
}  
  
group "default" {  
    targets = ["webapp"]  
}  
  
target "webapp" {  
    args = {  
        buildno = "${increment(123)}"  
    }  
}
```



Cache



Cache

Inline

Local

Registry

Amazon S3 (Experimental)

GitHub Action (Experimental)

Azure Blob Storage (Experimental)



Inline

Cache will be part of the image

Not scalable



Local

Cache inside a directory in your filesystem using OCI format

Good choice for testing



Registry

An extension of inline cache, but the cache will be stored as a separate image

Allows for separating the cache and resulting image artifacts so that you can distribute your final image without the cache inside.

It can efficiently cache multi-stage builds in max mode, instead of only the final stage.

It works with other exporters for more flexibility, instead of only the image exporter.



Exporters



Image and registry exporters

The image exporter outputs the build result into a container image format.

The registry exporter is identical, but it automatically pushes the result by setting `push=true`.



Local and tar exporters

The local and tar exporters output the root filesystem of the build result into a local directory.

They're useful for producing artifacts that aren't container images.



OCI and Docker exporters

The oci exporter outputs the build result into an OCI image layout tarball.

The docker exporter behaves the same way, except it exports a Docker image layout instead.



DEMO

